Machine learning glossary

Building blocks

- Machine learning (ML) in essence – approximating complex dependencies

Classical programming

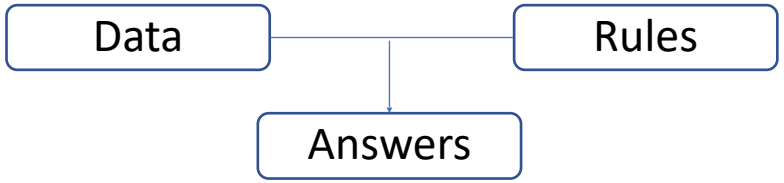

ML

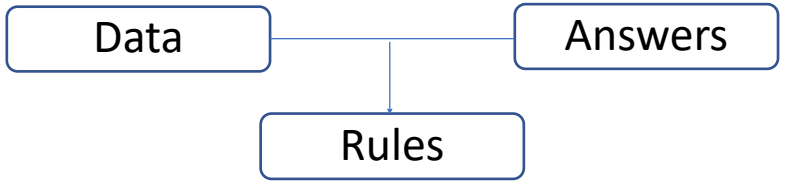

- Object – unit of input that gives us a distinct answer
  - Could be simple such as a client deal
  - Could be more complex such as recent 5 days of client activity
- Feature – numeric characteristics of an object
  - For example, how much net dv01 client X bought 3 days before
- Target = answer, this is what we would like to predict
  - Say, $RUB or OFZ yield movement

Zooming in on how data is represented for ML:

…

Model

Model – representation of rules

- Eg linear equation, or a chain of logical gates

Algorithm – procedure for efficient discovery of rules

- Closely related to model – often used as synonyms

Basic model assumption – nature of dependency:

- Non-linear
  - Most dependencies we deal with, especially flow-related ones
  - For example, if client X buys in combination with client Y, it might be more powerful than simply the sum of two
- Linear
  - Still, in cases where dependency is clearly linear (say, we model where RUB should be trading based on where MXN, ZAR etc are trading), a linear model will have advantage

Model output types:

- Regression – gives exact number
  - Eg size of move over 10 days horizon
- Classification – assigns a label
  - Say, whether first 1% move going forward will be up or down (binary classification)

Learning and evaluation

Data is split into train set and test set → Algorithm learns by minimizing error on train set → Quality is estimated on data it hasn't seen – the test set

Cross-validation technique – several different train/test splits for the same data, so that several estimates of quality are available

Metrics to assess quality

- Regression:
  - R2 – to which extent variance in the target variable is explained by the model
- Classification:
  - Accuracy – percentage of all objects we classified correctly (could be misleading)
  - Precision – out of objects we classified as 'up', how many were 'up' in reality
  - Recall – out of all objects that were in reality 'up', how many we correctly discovered as 'up'
  - ROC AUC – area under curve that shows tradeoff b/w TP and FP rates

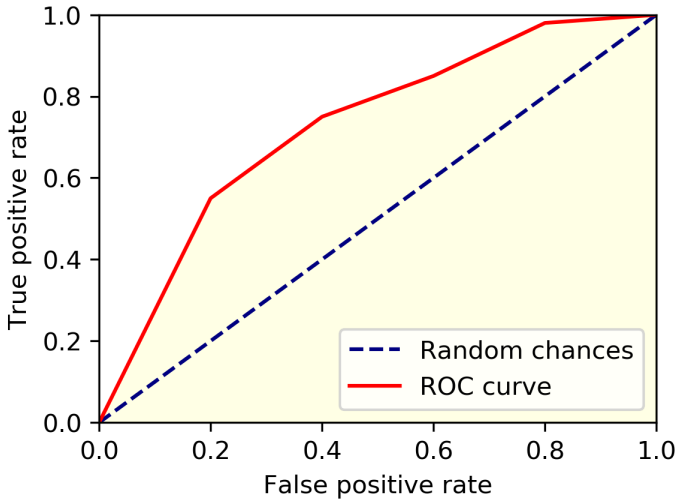

All of those can be calculated out of so-called confusion matrix, which plots labels that model put on objects against what they were in reality

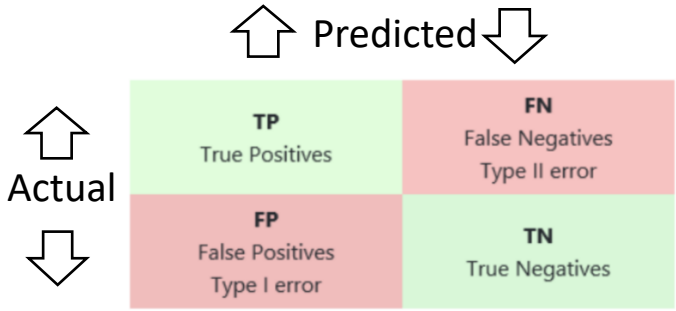

Decision tree ensemble (aka Random Forest)

Basic unit of the algorithm – decision tree

- nodes (blue) – logic gates
- leaves (orange) – predictions

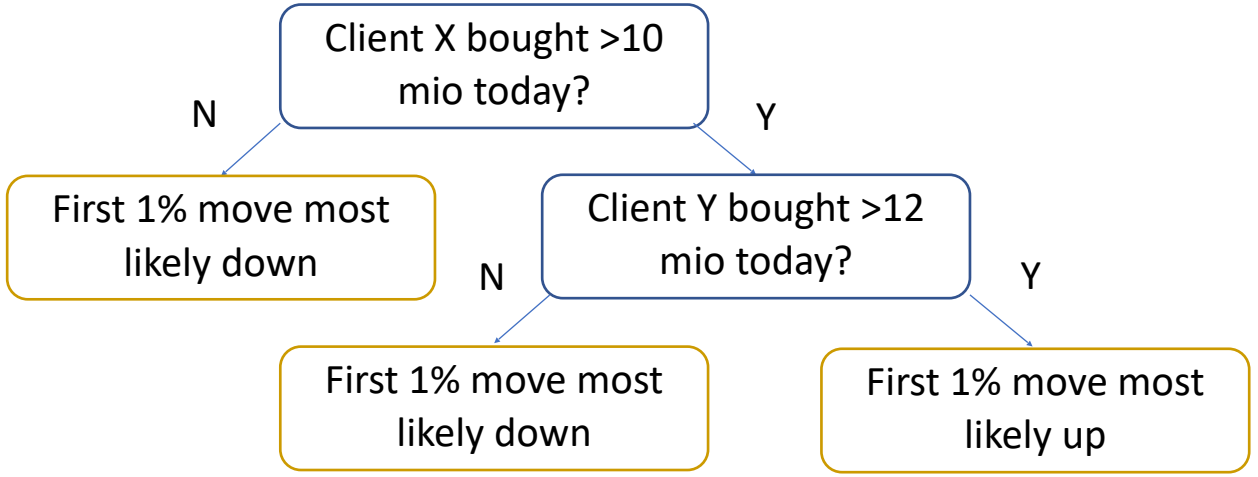

Decision tree goes over all features and chooses split that gives lowest entropy

- Entropy – measure of “chaos”, “randomness”

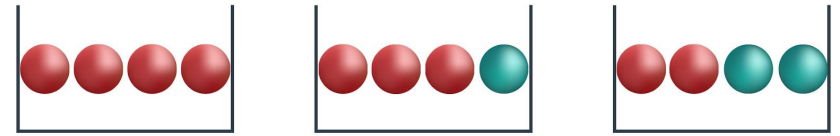

high knowledge
low entropy

medium knowledge
medium entropy

low knowledge
high entropy

Single tree vulnerability: unstable, change in input dataset leads to a different tree

Solution – ensembling (“forest” of trees):

- Eliminates instability by averaging predictions from many individual trees
- To further improve stability, each tree sees only a random subset of overall training data (method called bootstrap aggregation), so short name for the overall algorithm is Random Forest